JORDAN LEE

AI Platform Engineer | Retrieval, Evaluation, and Reliable LLM Systems

Fictional profile for the llm-client cookbook | jordan.lee@example.test | Toronto, Canada

SYNTHETIC EXAMPLE DOCUMENT	Safe for upload and retrieval demonstrations

PROFILE

Platform engineer focused on provider-agnostic LLM applications, typed tool execution, retrieval quality, and cost-aware production operations. Builds observable systems that expose model choice, request tier, cache behavior, and token-level usage without leaking provider-specific details into application code.

SELECTED EXPERIENCE

Northstar Systems | Senior AI Platform Engineer | 2023-present

- Designed a multi-provider inference gateway supporting OpenAI and Anthropic models, structured outputs, streaming, tool calls, retries, and explicit request-tier selection.
- Implemented prompt-cache observability and pricing attribution, including cached-input token accounting and provider-specific cache read/write rates.
- Built an evaluation harness for answer quality, tool correctness, latency, and cost; release candidates require deterministic offline checks plus opt-in live provider smoke tests.

Harbor Analytics | Machine Learning Engineer | 2020-2023

- Shipped hybrid retrieval using lexical search, embeddings, reranking, and citation-grounded response generation over technical documents.
- Reduced median retrieval latency by 34 percent through batched embedding requests, index filtering, and cache-aware query planning.
- Created failure-injection tests for rate limits, partial tool failures, malformed structured output, and provider timeouts.

FRONTIER CAPABILITY PROJECTS

Provider Batch Pipeline	Submitted asynchronous evaluation jobs through native OpenAI Batch and Anthropic Message Batches APIs, while preserving provider job identifiers and lifecycle status.
Typed Research Assistant	Combined JSON Schema outputs, namespaced tools, file search, and citation capture with bounded reasoning effort for predictable cost.
Cache-Aware Long-Context Review	Marked stable instructions and reference documents for provider caching, then reported cache reads, cache writes, execution mode, and estimated cost per request.

TECHNICAL TOOLKIT

LLM APIs	OpenAI Responses API, Anthropic Messages API, streaming, structured outputs, prompt caching, native tools, provider batch APIs
Runtime	Python, asyncio, Pydantic, HTTPX, retries, circuit breakers, idempotency, tracing
Retrieval	vector stores, embeddings, hybrid search, reranking, metadata filters, citations
Quality	pytest, contract tests, semantic benchmarks, drift checks, cost and latency budgets